

MongoDB Documentation

1. What is MongoDB?

MongoDB is a popular NoSQL database that stores data in a flexible, JSON-like format called **BSON**. Unlike relational databases, MongoDB doesn't use rows and columns, but stores data as documents inside collections. It is designed for **high performance**, **scalability**, and **ease of development**, making it ideal for modern web, real-time analytics, and big data applications.

Key Features

- **Document-Oriented:** Data is stored in collections of BSON documents, making it easy to model complex, hierarchical relationships.
- **Schema-Less:** Collections can hold documents with different fields, supporting evolving data requirements.
- **High Performance:** Fast read/write operations optimized for large-volume workloads.
- **Horizontal Scaling:** Integrated support for sharding and distributed architectures.
- **Rich Query Language:** Supports powerful filtering, sorting, aggregation, and indexing.

Example MongoDB Document

```
{  
  "name": "Ruthra",  
  "age": 21,  
  "city": "Avinashi"  
}
```

2. MongoDB Architecture

- **MongoDB Server (mongod):** Core engine handling data storage, access, and background tasks.
- **MongoDB Shell (mongosh):** Interactive JavaScript-based client to query and manage databases.
- **Database:** Logical container for collections.
- **Collection:** Group of BSON documents (like tables in SQL).
- **Document:** Atomic unit of data (like rows in SQL), stored as flexible BSON objects.

*Insight: MongoDB's architecture is developed to facilitate **distributed data storage** and management, with easy migration between deployment environments (local, cloud, on-premises).*

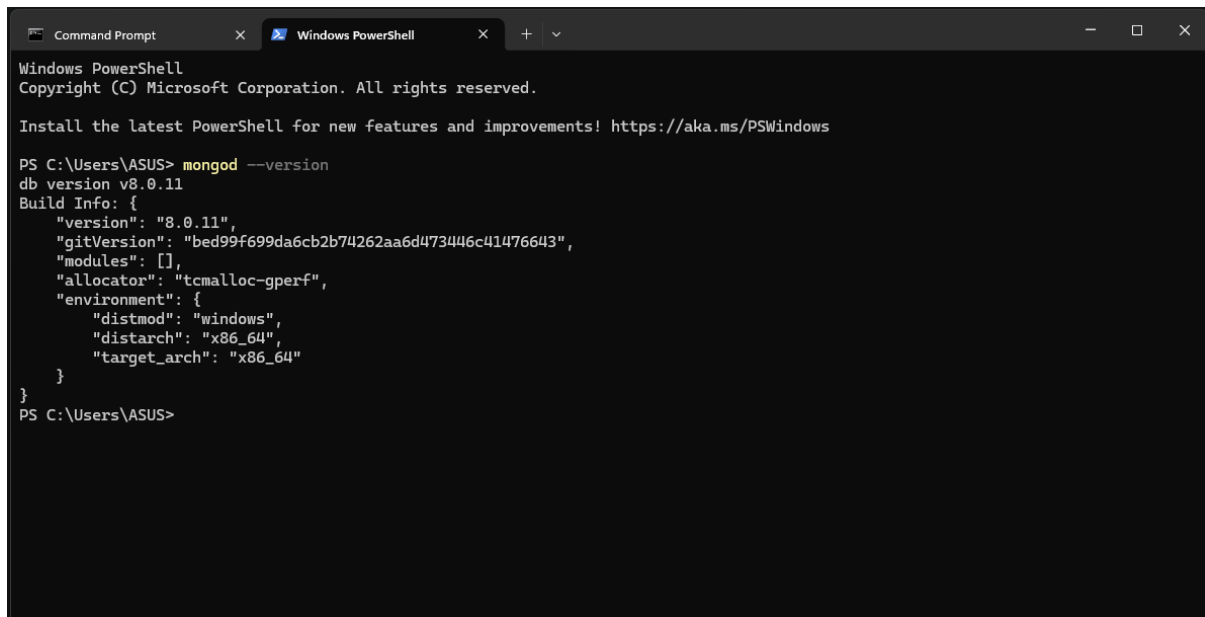
3. Installation

On Windows

1. Download MongoDB Community Edition from the [official site](#).
2. Choose the **MSI installer** for Windows.

3. Run the installer and select **Complete Setup**.
4. Enable **Install as a Service**—this allows MongoDB to start automatically with Windows.
5. Install **MongoDB Shell (mongosh)** (usually bundled; download separately if missing).
6. (Optional) Install **MongoDB Compass**, a GUI tool for database interaction.
7. **Verify** installation by running mongosh in Command Prompt: this should launch the MongoDB shell.

Verify Installation:



```

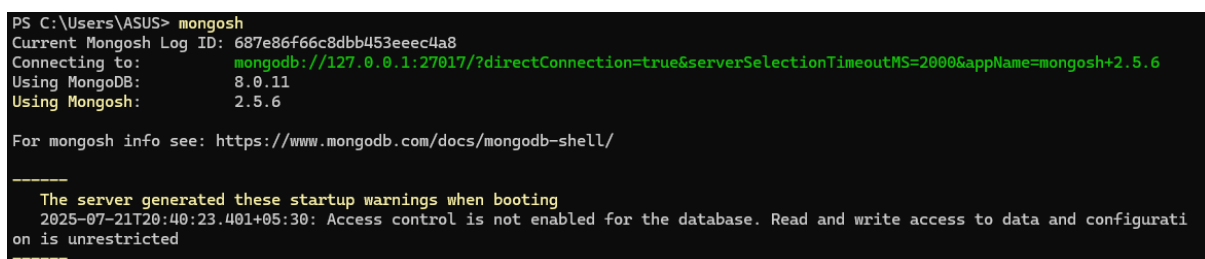
Windows PowerShell
Copyright (C) Microsoft Corporation. All rights reserved.

Install the latest PowerShell for new features and improvements! https://aka.ms/PSWindows

PS C:\Users\ASUS> mongod --version
db version v8.0.11
Build Info: {
  "version": "8.0.11",
  "gitVersion": "bed99f699da6cb2b74262aa6d473446c41476643",
  "modules": [],
  "allocator": "tcmalloc-gperf",
  "environment": {
    "distmod": "windows",
    "distarch": "x86_64",
    "target_arch": "x86_64"
  }
}
PS C:\Users\ASUS>

```

To Start the MongoDB Server use command ‘mongosh’



```

PS C:\Users\ASUS> mongosh
Current Mongosh Log ID: 687e86f66c8dbb453eeec4a8
Connecting to:      mongodb://127.0.0.1:27017/?directConnection=true&serverSelectionTimeoutMS=2000&appName=mongosh+2.5.6
Using MongoDB:      8.0.11
Using Mongosh:      2.5.6

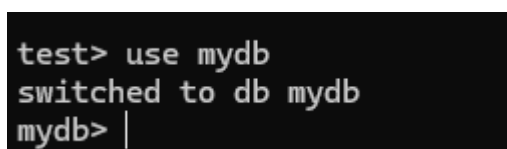
For mongosh info see: https://www.mongodb.com/docs/mongodb-shell/

-----
  The server generated these startup warnings when booting
  2025-07-21T20:40:23.401+05:30: Access control is not enabled for the database. Read and write access to data and configurati
on is unrestricted
-----

```

4. Core MongoDB Operations

Creating and Switching Databases



```

test> use mydb
switched to db mydb
mydb> |

```

Switches to (or creates) the database "mydb".

Creating Collections and Inserting Documents

```

mydb> db.createCollection("people")
{ ok: 1 }
mydb> db.people.insertOne({ name: "Alice", age: 25, city: "Chennai" })
{
  acknowledged: true,
  insertedId: ObjectId('687e87936c8dbb453eeec4a9')
}
mydb> db.people.insertMany([
...   { user_id: 1, age: 25, status: "passed" },
...   { user_id: 2, age: 30, status: "failed" }
... ])
...
{
  acknowledged: true,
  insertedIds: {
    '0': ObjectId('687e87a86c8dbb453eeec4aa'),
    '1': ObjectId('687e87a86c8dbb453eeec4ab')
  }
}
mydb> |

```

Viewing Databases and Collections

```

mydb> show dbs
admin          40.00 KiB
config         48.00 KiB
hexaware       24.00 KiB
local          112.00 KiB
mydb           40.00 KiB
nvd_database   12.00 KiB
school         100.00 KiB
mydb> show collections
people
mydb> |

```

Fetching Records

```

mydb> db.people.find()
[
  {
    _id: ObjectId('687e87936c8dbb453eeec4a9'),
    name: 'Alice',
    age: 25,
    city: 'Chennai'
  },
  {
    _id: ObjectId('687e87a86c8dbb453eeec4aa'),
    user_id: 1,
    age: 25,
    status: 'passed'
  },
  {
    _id: ObjectId('687e87a86c8dbb453eeec4ab'),
    user_id: 2,
    age: 30,
    status: 'failed'
  }
]
mydb> |

```

```

mydb> db.people.find({}, { user_id: 1, status: 1 }) // Specific fields
[
  { _id: ObjectId('687e87936c8dbb453eeec4a9') },
  {
    _id: ObjectId('687e87a86c8dbb453eeec4aa'),
    user_id: 1,
    status: 'passed'
  },
  {
    _id: ObjectId('687e87a86c8dbb453eeec4ab'),
    user_id: 2,
    status: 'failed'
  }
]
mydb> |

mydb> db.people.find({}, { user_id: 1, status: 1, _id: 0 }) // Specific fields, hide _id
[
  {},
  { user_id: 1, status: 'passed' },
  { user_id: 2, status: 'failed' }
]
mydb> |

```

```

mydb> db.people.findOne() // Single document
{
  _id: ObjectId('687e87936c8dbb453eeec4a9'),
  name: 'Alice',
  age: 25,
  city: 'Chennai'
}
mydb> |

```

Result is shown in an array of JSON-like documents.

Querying Data (Filter, Logical)

```

mydb> db.people.find({ status: "passed" }) // Filter by value
[
  {
    _id: ObjectId('687e87a86c8dbb453eeec4aa'),
    user_id: 1,
    age: 25,
    status: 'passed'
  }
]
mydb> db.people.find({ age: { $gt: 25 } }) // Greater than
[
  {
    _id: ObjectId('687e87a86c8dbb453eeec4ab'),
    user_id: 2,
    age: 30,
    status: 'failed'
  }
]
mydb> db.people.find({ $or: [{ status: "passed" }, { age: 25 }] }) // OR condition
[
  {
    _id: ObjectId('687e87936c8dbb453eeec4a9'),
    name: 'Alice',
    age: 25,
    city: 'Chennai'
  },
  {
    _id: ObjectId('687e87a86c8dbb453eeec4aa'),
    user_id: 1,
    age: 25,
    status: 'passed'
  }
]
mydb> db.people.find({ status: "passed", age: 25 }) // AND condition
[
  {
    _id: ObjectId('687e87a86c8dbb453eeec4aa'),
    user_id: 1,
    age: 25,
    status: 'passed'
  }
]

```

Sorting and Limiting Results

```
mydb> db.people.find().sort({ user_id: 1 }) // Ascending
[
  {
    _id: ObjectId('687e87936c8dbb453eeec4a9'),
    name: 'Alice',
    age: 25,
    city: 'Chennai'
  },
  {
    _id: ObjectId('687e87a86c8dbb453eeec4aa'),
    user_id: 1,
    age: 25,
    status: 'passed'
  },
  {
    _id: ObjectId('687e87a86c8dbb453eeec4ab'),
    user_id: 2,
    age: 30,
    status: 'failed'
  }
]
mydb> db.people.find().sort({ user_id: -1 }) // Descending
[
  {
    _id: ObjectId('687e87a86c8dbb453eeec4ab'),
    user_id: 2,
    age: 30,
    status: 'failed'
  },
  {
    _id: ObjectId('687e87a86c8dbb453eeec4aa'),
    user_id: 1,
    age: 25,
    status: 'passed'
  },
  {
    _id: ObjectId('687e87936c8dbb453eeec4a9'),
    name: 'Alice',
    age: 25,
    city: 'Chennai'
  }
]
mydb> |
```

`db.people.find().limit(5).skip(10)` // Pagination: skipping and limiting

Counting and Distinct

```
mydb> db.people.find().limit(5).skip(10) // Pagination: skipping and limiting
mydb> db.people.countDocuments({}) // Total records
3
mydb> db.people.countDocuments({ age: { $gt: 30 } }) // With condition
0
mydb> db.people.distinct("status") // All unique status values
[ 'failed', 'passed' ]
mydb> |
```

Counting helps in analytics and reporting; distinct is used for categorical breakdowns.

Updating and Deleting

```
[ 'failed', 'passed' ]
mydb> db.people.updateOne({ user_id: 1 }, { $set: { age: 26 } }) // Update record
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
mydb> db.people.deleteOne({ user_id: 1 }) // Delete record
{ acknowledged: true, deletedCount: 1 }
mydb> |
```

Dropping Collections and Databases

```
mydb> db.people.drop()      // Remove the collection
true
mydb> db.dropDatabase()     // Remove the whole database
{ ok: 1, dropped: 'mydb' }
mydb> |
```

Explain Query Plans (Performance Insight)

```
mydb> db.people.find({ status: "A" }).explain()
{
  explainVersion: '1',
  queryPlanner: {
    namespace: 'mydb.people',
    parsedQuery: { status: { '$eq': 'A' } },
    indexFilterSet: false,
    optimizationTimeMillis: 0,
    maxIndexedOrSolutionsReached: false,
    maxIndexedAndSolutionsReached: false,
    maxScansToExplodeReached: false,
    prunedSimilarIndexes: false,
    winningPlan: { isCached: false, stage: 'EOF' },
    rejectedPlans: []
  },
  queryShapeHash: 'E4871C1FC6A4544FA07E582BF9C76628598DFC2DA2B46B02BC8B86D1A595E8BD',
  command: { find: 'people', filter: { status: 'A' }, '$db': 'mydb' },
  serverInfo: {
    host: 'Ruthra-PC',
    port: 27017,
    version: '8.0.11',
    gitVersion: 'bed99f699da6cb2b74262aa6d473446c41476643'
  },
  serverParameters: {
    internalQueryFacetBufferSizeBytes: 104857600,
    internalQueryFacetMaxOutputDocSizeBytes: 104857600,
    internalLookupStageIntermediateDocumentMaxSizeBytes: 104857600,
    internalDocumentSourceGroupMaxMemoryBytes: 104857600,
    internalQueryMaxBlockingSortMemoryUsageBytes: 104857600,
    internalQueryProhibitBlockingMergeOnMongoS: 0,
    internalQueryMaxAddToSetBytes: 104857600,
    internalDocumentSourceSetWindowFieldsMaxMemoryBytes: 104857600,
    internalQueryFrameworkControl: 'trySbeRestricted',
    internalQueryPlannerIgnoreIndexWithCollationForRegex: 1
  },
  ok: 1
}
mydb> |
```

This shows how MongoDB will execute the query, useful for optimizing performance in larger databases.

5. Sample CRUD Workflow in MongoDB

1. **Connect to MongoDB:** Ensure mongod is running, then open mongosh.
2. **Create/Choose Database:** use demodb
3. **Create Collection:** db.createCollection("people")
4. **Insert Data:** Use insertOne or insertMany.
5. **Read Data:** Use find, findOne, with optional filters and projections.
6. **Update Data:** Use updateOne, \$set for field updates.
7. **Delete Data:** Use deleteOne or deleteMany.
8. **Drop Collection/Database:** Use drop/dropDatabase as needed.

6. Practical Example

Suppose you have a database called trainingdb and a collection people:

```
mydb> db.people.insertMany([
...   { user_id: "bc101", status: "A", age: 25 },
...   { user_id: "bc102", status: "B", age: 30 },
...   { user_id: "bc103", status: "A", age: 50 },
...   { user_id: "xy201", status: "A", age: 35 },
...   { user_id: "mn301", status: "B", age: 20 }
... ])
...
{
  acknowledged: true,
  insertedIds: {
    '0': ObjectId('687e899e6c8dbb453eeec4ac'),
    '1': ObjectId('687e899e6c8dbb453eeec4ad'),
    '2': ObjectId('687e899e6c8dbb453eeec4ae'),
    '3': ObjectId('687e899e6c8dbb453eeec4af'),
    '4': ObjectId('687e899e6c8dbb453eeec4b0')
  }
}
mydb> |
```

- Find all: db.people.find()

```
mydb> db.people.find()
[
  {
    _id: ObjectId('687e899e6c8dbb453eeec4ac'),
    user_id: 'bc101',
    status: 'A',
    age: 25
  },
  {
    _id: ObjectId('687e899e6c8dbb453eeec4ad'),
    user_id: 'bc102',
    status: 'B',
    age: 30
  },
  {
    _id: ObjectId('687e899e6c8dbb453eeec4ae'),
    user_id: 'bc103',
    status: 'A',
    age: 50
  },
  {
    _id: ObjectId('687e899e6c8dbb453eeec4af'),
    user_id: 'xy201',
    status: 'A',
    age: 35
  },
  {
    _id: ObjectId('687e899e6c8dbb453eeec4b0'),
    user_id: 'mn301',
    status: 'B',
    age: 20
  }
]
mydb> |
```

- Filter: `db.people.find({ status: "A" })`

```
mydb> db.people.find({ status: "A" })
[
  {
    _id: ObjectId('687e899e6c8dbb453eeec4ac'),
    user_id: 'bc101',
    status: 'A',
    age: 25
  },
  {
    _id: ObjectId('687e899e6c8dbb453eeec4ae'),
    user_id: 'bc103',
    status: 'A',
    age: 50
  },
  {
    _id: ObjectId('687e899e6c8dbb453eeec4af'),
    user_id: 'xy201',
    status: 'A',
    age: 35
  }
]
mydb> |
```

- Sort: `db.people.find().sort({ user_id: 1 })`

```
mydb> db.people.find().sort({ user_id: 1 })
[
  {
    _id: ObjectId('687e899e6c8dbb453eeec4ac'),
    user_id: 'bc101',
    status: 'A',
    age: 25
  },
  {
    _id: ObjectId('687e899e6c8dbb453eeec4ad'),
    user_id: 'bc102',
    status: 'B',
    age: 30
  },
  {
    _id: ObjectId('687e899e6c8dbb453eeec4ae'),
    user_id: 'bc103',
    status: 'A',
    age: 50
  },
  {
    _id: ObjectId('687e899e6c8dbb453eeec4b0'),
    user_id: 'mn301',
    status: 'B',
    age: 20
  },
  {
    _id: ObjectId('687e899e6c8dbb453eeec4af'),
    user_id: 'xy201',
    status: 'A',
    age: 35
  }
]
mydb> |
```

- Count: `db.people.countDocuments({})`

```
mydb> db.people.countDocuments({})
5
mydb> |
```


- Pattern: `db.people.find({ user_id: /^bc/ })`

```
mydb> db.people.find({ user_id: /^bc/ })
[
  {
    _id: ObjectId('687e899e6c8dbb453eeec4ac'),
    user_id: 'bc101',
    status: 'A',
    age: 25
  },
  {
    _id: ObjectId('687e899e6c8dbb453eeec4ad'),
    user_id: 'bc102',
    status: 'B',
    age: 30
  },
  {
    _id: ObjectId('687e899e6c8dbb453eeec4ae'),
    user_id: 'bc103',
    status: 'A',
    age: 50
  }
]
mydb> |
```

- Distinct: `db.people.distinct("status")`

```
mydb> db.people.distinct("status")
[ 'A', 'B' ]
mydb> |
```

- Update: `db.people.updateOne({ user_id: "bc101" }, { $set: { age: 26 } })`

```
mydb> db.people.updateOne({ user_id: "bc101" }, { $set: { age: 26 } })
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
mydb> |
```

- Delete: `db.people.deleteOne({ user_id: "mn301" })`

```
mydb> db.people.deleteOne({ user_id: "mn301" })
{ acknowledged: true, deletedCount: 1 }
mydb> |
```

This covers all fundamental MongoDB CRUD operations you'll commonly use in real-world projects.

7. Additional Expert Insights

- **Schema Design:** Though MongoDB is schema-less, plan document structures up front for consistency and efficiency, using embedded documents for “belongs-to” scenarios and references for one-to-many relationships in large datasets.
- **Indexing:** Create indexes on frequently queried fields (`db.collection.createIndex()`) to significantly improve query performance.
- **Sharding:** For very large datasets, consider sharding collections to distribute data across multiple servers and balance load.
- **Aggregation Framework:** Use `db.collection.aggregate()` for complex data analysis—grouping, summarizing, and transforming data.
- **Security:** Enable authentication, role-based access, and network-level protections in production deployments.
- **Backup and Restore:** Use `mongodump` and `mongorestore` for efficient database backup and recovery.
- **Monitoring:** Use MongoDB’s built-in tools or external solutions (e.g., Ops Manager, Atlas) to monitor performance and get alerts on anomalies.

8. Conclusion

MongoDB’s flexible architecture, rich query language and scalability features make it a top choice for developers building data-intensive applications. Mastering the foundational commands, patterns, and best-practices outlined above sets a strong base for both learning and deploying full-scale MongoDB solutions.